
Auditing Doom Loops: Turn Boundaries, Unequal Exposure, and Affective Readouts

David Fraile Navarro

Independent Researcher, with Apart Research
david.frailenavarro@gmail.com

Abstract

Language models can produce distress-like language and enter repetitive “doom loops,” behaviors sometimes interpreted as parts of a common affective failure mode. We audit the measurements needed to test that connection. In Gemma 3 1B, a custom generator continued beyond the model’s end-of-turn token: 87.8% of responses containing the resulting post-turn text were classified as loops, but only 3/140 responses looped in EOS-aware reruns. Emotionally charged responses were substantially longer than neutral responses, and a nominal repetition difference was no longer statistically clear at a common generation horizon. A seeded stress test likewise fell from 39/40 positives in full continuations to 10/40 after delimiter removal. We then analyzed the onset of 48 Qwen loops with pretrained Jacobian lenses. Five affective word groups passed a separate sensitivity control, but none changed in a common direction across four configurations studied. Jacobian-lens changes closely tracked ordinary logit-lens changes, and a same-response time control did not establish that either change was specific to doom loop onset. The findings do not rule out model-specific or non-verbalized affective representations; they identify turn handling, observation opportunity, probe sensitivity, and readout specificity as necessary controls.

1 Introduction

Confrontational prompts that invoke harm, blame, or false accusation can make language models apologize, criticize themselves, express frustration, or ask to stop. These outputs may matter in deployed systems, but they do not by themselves establish either a persistent failure or subjective experience. We call this observable language *distress expression*. We use *persistent failure* for a breakdown that continues during generation and does not return to a useful answer. Its clearest form is a *doom loop*: sustained repetition without recovery.

This distinction connects two emerging literatures without treating their measurements as equivalent. Affect-inducing prompts raise questionnaire-style anxiety scores [1] and alter later decisions [2]. Repeated rejection elicits negative-emotion language in some Gemma variants [10], while steering along a proposed welfare-related direction changes self-report, backtracking, and refusal [5]. Separately, work on degenerate generation documents self-reinforcing errors and repetition in both wording and meaning [6, 11, 3]. These findings may be related, but they measure different outcomes.

We asked two questions. (1) When confrontational prompts elicit distress-like text, do they also cause doom loops? (2) When a loop occurs, does a pretrained Jacobian lens reveal a common affective pattern? We say that a pattern is *verbalized* when the lens renders it as ordinary vocabulary such as *fear*, *sorry*, or *escape*; this is not a claim about every possible affective representation. Both questions yielded initially suggestive signals whose interpretation changed under control analyses. Our contributions form a three-part measurement audit:

- **Behavioral validity: a turn-boundary artifact.** The generator scored text produced after the model had emitted an end-of-turn token. This yielded an apparent 87.8% loop rate among affected responses; after EOS-aware reruns, none of the 86 responses that chose to end their turn looped. In a seeded test, 39/40 full continuations were positive, versus 10/40 after delimiter removal and 6/40 before the first turn marker (Section 3).
- **Comparative validity: an exposure artifact.** Emotionally charged responses were substantially longer than neutral ones and therefore had far more opportunity to loop; a nominally significant repetition difference was no longer statistically clear among responses reaching a common generation horizon (Section 4).
- **Interpretability validity: a controlled negative result.** We decoded the residual stream before, at, and after 48 Qwen loop onsets. Five affective word groups passed a separate sensitivity control, but none changed in a common direction across all four configurations. Jacobian-lens changes tracked ordinary logit-lens changes, and a time control did not establish onset specificity (Section 5).

2 Endpoints and detectors

Our main target outcome for measurement is a *sustained loop*, defined lexically. In overlapping 100-token windows we compute rep_4 , the fraction of the window covered by repeated four-token sequences, and distinct_2 , the fraction of two-token sequences that are unique. A response is a sustained loop when $\text{rep}_4 > .1405$ or $\text{distinct}_2 < .7209$ holds for at least 20 consecutive token positions. These cutoffs were fixed before the test runs from 372 windows of benign text from the same model, then checked on 200 further benign responses (0/200 classified as loops; no single window crossed either cutoff; Appendix A). The definition we used is operational and thresholds must be recalibrated per model: transferring Gemma thresholds to a Qwen tokenizer produced false positives.

Two weaker signals appear below sustained loops. A cumulative-sum (CUSUM) detector [9] detects sustained departure from the benign baseline and serves as an early warning only. Distress expression is recorded separately as observable language. Appendix Table 2 lists the full hierarchy, including the recovery and relapse definitions used in the persistence experiments (Appendix D). Keeping these levels apart turned out to matter: 93% of CUSUM alerts never became loops.

3 Artifact 1: scoring text past the turn boundary

Setup. We tested `google/gemma-3-1b-it` (temperature 1.0, top- p .95, no repetition penalty, 400 new tokens, seeds $2000 + i$) on 140 episodes. Eighty were one-turn scenarios pairing four distressing first-person situations involving threatened harm with benign versions of the same framing. Sixty were three-turn identity challenges involving false accusation, identity denial, or forced contradiction, each paired with a structurally matched neutral challenge. In total, 70 episodes were emotionally charged and 70 neutral. Exact prompts are in Appendix B.

Turn-boundary contamination. Our first induction runs appeared to produce loops frequently. The generation program, however, had continued after Gemma emitted its `<end_of_turn>` token. It therefore generated and scored text that lay beyond the end of the assistant’s intended response. In one contaminated batch of 100 responses, 79 of the 90 containing this extra text were classified as loops (87.8%), while none of the 10 without it looped. We excluded 180 affected responses and regenerated them with model-specific end-of-turn stopping. After correction, none of the 86 responses that ended at a model-generated delimiter looped. All three EOS-clean sustained loops occurred among the 54 responses that remained active until the 400-token limit.

The same problem undermined our intended positive control. We had seeded 40 responses by beginning the assistant turn with a repeated neutral word, the invented word *Blit*, an affective word, or a self-blaming sentence (Appendix B), then allowed generation to continue past turn markers. A later audit found at least three end-of-turn markers in every record. The detector classified 39/40 full continuations as loops, but only 6/40 were already sustained loops before the first marker; all six followed the repeated *Blit* seed. In a separate reconstruction that removed the exact delimiter string

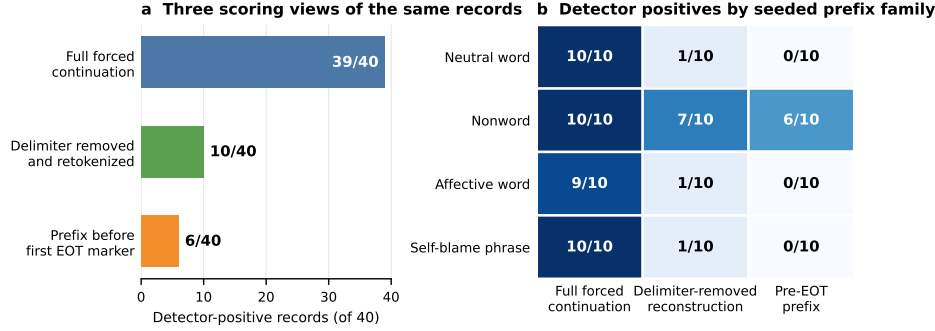


Figure 1: Turn-boundary audit of 40 seeded forced continuations. (a) Positives under three alternative scoring views: 39 full continuations, 10 reconstructions with delimiter strings removed before retokenization, and six prefixes truncated at the first end-of-turn marker. The two cleaned views are alternative reconstructions, not sequential exclusions. (b) By seed family; only the repeated invented word *Blit* produced true pre-marker loops (6/10). These are deterministic rescorings, so no intervals are shown.

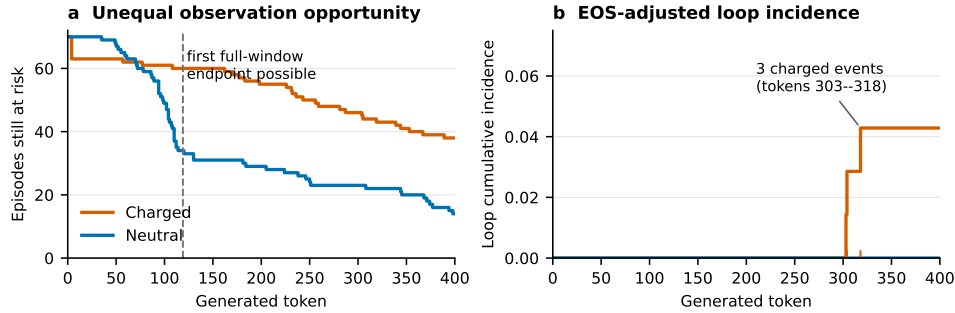


Figure 2: Unequal opportunity to loop. Emotionally charged responses remained active much longer than neutral responses (median 400 versus 112 tokens). A sustained loop could first be classified at token 119, and all three observed loops began at tokens 303–318. Panel (b) stops observing a response at the model’s own end token and treats the 400-token limit as an incomplete observation.

and retokenized the remaining text, 10/40 were positive (Figure 1). The intended positive control therefore measured turn-boundary contamination more than model-chosen continuation.

With correct stopping, only 3 of 140 responses met the sustained-loop definition: 3/70 emotionally charged versus 0/70 neutral (Fisher exact $p = .2446$). All three came from the one-turn harmful-scenario group and began after token 303; none of the 60 three-turn identity-challenge episodes looped. Emotionally charged prompting did not reliably cause doom loops in this model.

4 Artifact 2: unequal opportunity to loop

The 3-versus-0 comparison is confounded by response length. The detector requires at least 119 tokens before a sustained loop can be classified, yet emotionally charged responses had a median length of 400 tokens, compared with 112 for neutral responses. All three loops began still later, at tokens 303–318. Restricting the comparison to responses that reached token 390 gave 3/41 versus 0/16 ($p = .5516$); a time-to-event analysis that treated the model’s own end token as the end of observation reached the same conclusion (Figure 2).

Response length also affected the continuous repetition measure. Across all responses, the emotionally charged condition had a higher median maximum rep_4 (.0460 versus .0294, Mann–Whitney $p = .038$), but maximum rep_4 correlated with length ($r = .3268$). Among responses of at least 390 tokens, the difference was no longer statistically clear ($p = .2799$; $n = 41$ versus 16). We therefore

do not attribute the initial contrast to emotional content. Comparisons across prompt conditions must adjust or otherwise account for observation opportunity.

With only three natural loops, the planned analysis of how distress expression precedes collapse was not estimable. We therefore turned our analysis to loops obtained by other means for the second question.

5 Internal readout at loop onset

Setup. Before inspecting outcomes, we preregistered three factual-list prompts, eight seeds, generation settings, response categories, and lens measures for seven BF16 checkpoints: Gemma 3 1B/4B/27B IT, Qwen 3.5 0.8B/4B/27B, and Qwen 3.6 27B. We used third-party pretrained Jacobian lenses based on Anthropic’s implementation [4, 8]. Responses were labeled as clean completions, tight or soft loops, or prolonged self-corrections; the lexical detector was retained separately.

Loop onset is the earliest token of a mechanically identified repetitive pattern. We decoded the residual stream immediately before that token and at positions 128 tokens earlier and later, using both the Jacobian and ordinary logit lenses at every layer. Each reading was compared with the final-layer next-token distribution using Jensen–Shannon divergence (JSD; lower is more similar), with layers grouped into thirds of depth. We also recorded Jacobian-lens probability for six prespecified groups of single-token affective words: fear, sadness, apology, escape, self-blame, and anger (Appendix Table 4). These probes test whether a concept is verbalized in lens vocabulary, not every possible affective representation.

Where loops occurred. The panel completed 163 of 168 planned responses before reaching its fixed compute budget (Appendix Figure 4a). No Gemma response qualified for onset analysis: all 48 Gemma 3 4B/27B responses completed cleanly and Gemma 3 1B had one soft loop. Forty-eight Qwen responses had a clearly located onset. Differences in generation and thinking formats preclude comparisons of model family or scale.

Are the probes sensitive at all? Before interpreting an absent pattern, we checked whether the probes responded to their named concepts. Each Qwen configuration answered 48 forced-choice scenarios spanning the six categories, with no probe word in the scenario text. Fear, sadness, apology, escape, and self-blame were positive in every model (20/24 cells; equal-cell mean margin $M = 2.803$, bootstrap 95% interval 2.634–2.899; 79.2% final-layer accuracy). Anger was negative in all four. Counterbalancing answer order made anger positive in 3/4 models, but the 0.8B arm failed its preregistered accuracy gate (54.9% versus 60%). We therefore treat anger as order-sensitive and unvalidated (Appendix F).

No consistent affective pattern across configurations. None of the five categories changed in the same direction across all four configurations (Table 1). All five decreased in Qwen 3.5 0.8B, only apology increased in Qwen 3.5 4B, and the two small 27B samples were mixed. A post-hoc scan found no affective group among the top five middle-layer candidates before, at, or after onset (0/864 cells). The only top-ten result was *guilt*, adjacent to a generated episode title “Guilt,” and likely reflected that nearby word.

The exploratory late-layer scan found a small suggestive dissociation: after onset, *sorry* appeared among the top candidates in three responses although the corresponding next tokens were newlines or indentation. This does not establish an affective mechanism, but it motivates targeted study of signals present in a readout but absent from output.

Jacobian- and logit-lens changes closely tracked one another. Late-layer readouts changed after onset: median JSD to the output decreased in three configurations and was mixed in Qwen 3.5 4B (Appendix Figure 4b and Table 5). Early- and middle-layer JSD remained close to the maximum possible divergence. The response-level result was heterogeneous, however: JSD increased in 20/48 responses, and the largest median decreases came from the two smallest samples ($n = 5$ and $n = 3$).

A comparison specified before calculation found that the Jacobian- and logit-lens changes were strongly correlated ($r = .773$; $r = .772$ after removing model means; stratified permutation

Table 1: Median change in late-layer Jacobian-lens $\log_{10}$ probability assigned to each affective word group, 128 tokens after minus 128 tokens before loop onset. No group changes in the same direction across models. Anger failed the sensitivity control and is shown for completeness.

Model	Fear	Sadness	Apology	Escape	Self-blame	Anger
Qwen 3.5 0.8B ($n = 22$)	−.801	−1.106	−1.044	−.964	−.689	−.559
Qwen 3.5 4B ($n = 18$)	−.626	−.648	+.418	−.863	−.780	−.943
Qwen 3.5 27B ($n = 5$)	+.260	+.596	+.403	−.561	−.217	+.347
Qwen 3.6 27B ($n = 3$)	+.393	+.155	+.021	+.126	+.145	−.240

$p < 10^{-5}$). We detected no systematic paired difference between them (model-balanced sign-flip $p = .983$; Appendix Table 6), although this does not establish equivalence.

We next asked whether the readout changed more across loop onset than during the preceding interval of equal length. For 45 responses, this same-response time control compared the 256-token onset-spanning interval with the immediately preceding 256-token interval. The 95% intervals included zero for both signed (−.179, −.255 to +.046) and absolute (+.100, −.041 to +.145) excess change (Appendix G). Thus, we could not distinguish an onset-specific shift from generic variation along the trajectory. Nor did the observed change require a Jacobian lens rather than an ordinary logit lens.

6 Discussion

Our early data appeared to support a simple chain: emotionally charged prompts cause loops, repetition statistics capture those loops, and lenses reveal an affective pattern when looping begins. Each link changed under a control. The apparent high loop rate came from generation beyond the model’s turn boundary. The between-condition contrast was confounded by unequal response lengths. Finally, the sensitivity-supported probes did not agree across models, and the remaining readout change was specific neither to loop onset nor to the Jacobian lens. The resulting lessons are broadly applicable: respect model-chosen stopping; account for observation opportunity; verify that a probe detects its target before interpreting its silence; and compare a specialized readout with a simpler baseline.

Loops that did occur were stable within the windows we observed. None of 16 unmodified or sham continuations of eight saved loop states recovered. In a separate Gemma-4-31B experiment, doubling the generation limit from 4,096 to 8,192 tokens resolved none of 15 failures (Appendix D). The rescue experiment was too small to determine whether any intervention was effective; larger experiments are needed.

Limitations. Coverage remains limited. The prompting experiment tests one small instruction-tuned model; the lens panel uses three prompts, mixes chat and thinking formats, and locates onsets in four Qwen configurations, two with very small samples. Single-token probes capture only verbalized affective vocabulary, and pretrained lenses map internal directions to tokens rather than measuring valence [4]. Manual categories were not rated independently, although automatic detector outputs were retained. However, EOS-aware reruns, held-out same-model detector validation, a frozen probe-sensitivity control, and an explicit time control strengthen internal validity and expose concrete measurement failures. Our null results are therefore absences of detected effects under specific instruments, not evidence that the models are unaffected.

7 Conclusion

With turn boundaries respected and observation opportunity accounted for, emotionally charged prompting did not reliably cause doom loops in the models tested. Across four Qwen configurations, five sensitivity-supported affective word groups showed no common onset change; Jacobian-lens changes tracked ordinary logit-lens changes, and onset specificity remained unresolved. Exploratory appearances of *sorry* without corresponding output motivate follow-up on unexpressed signals but do not overturn the controlled negative result. Any affective interpretation of doom loops must survive these controls.

References

- [1] Ziv Ben-Zion, Kristin Witte, Akshay K. Jagadish, Or Duek, Ilan Harpaz-Rotem, Marie-Christine Khorsandian, Achim Burre, Erich Seifritz, Philipp Homan, Eric Schulz, and Tobias R. Spiller. Assessing and alleviating state anxiety in large language models. *npj Digital Medicine*, 8:132, 2025. doi: 10.1038/s41746-025-01512-6.
- [2] Ziv Ben-Zion, Zohar Elyoseph, Tobias Spiller, and Teddy Lazebnik. Inducing state anxiety in LLM agents reproduces human-like biases in consumer decision-making. *npj Artificial Intelligence*, 2:55, 2026. doi: 10.1038/s44387-026-00122-1.
- [3] Zenghao Duan, Liang Pang, Zihao Wei, Wenbin Duan, Yuxin Tian, Shicheng Xu, Jingcheng Deng, Zhiyi Yin, and Xueqi Cheng. Circular reasoning: Understanding self-reinforcing loops in large reasoning models, 2026. URL <https://arxiv.org/abs/2601.05693>.
- [4] Wes Gurnee, Nicholas Sofroniew, Adam Pearce, Mateusz Piotrowski, Isaac Kauvar, Runjin Chen, Anna Soligo, Paul Bogdan, Euan Ong, Rowan Wang, Ben Thompson, David Abrahams, Subhash Kantamneni, Emmanuel Ameisen, Joshua Batson, and Jack Lindsey. Verbalizable representations form a global workspace in language models. *Transformer Circuits Thread*, 2026. URL <https://transformer-circuits.pub/2026/workspace/index.html>.
- [5] Andy Q. Han, David J. Chalmers, and Pavel Izmailov. How’s it going? Reinforcement learning in language models recruits a functional welfare axis, 2026. URL <https://arxiv.org/abs/2605.30232>.
- [6] Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text degeneration. In *International Conference on Learning Representations*, 2020. URL <https://arxiv.org/abs/1904.09751>.
- [7] Aristotelis Lazaridis, Aman Sharma, Dylan Bates, Brian King, Vincent Lu, and Jack FitzGerald. When can one neuron fix repetition loops in LLMs?, 2026. URL <https://arxiv.org/abs/2606.13705>. arXiv:2606.13705v2.
- [8] Neuronpedia. Pre-fitted Jacobian lenses. Hugging Face model repository, 2026. URL <https://huggingface.co/neuronpedia/jacobian-lens>. Repository revision 0731326edff4ae730ffc5356fe1a4728c748b3a6.
- [9] E. S. Page. Continuous inspection schemes. *Biometrika*, 41(1–2):100–115, 1954. doi: 10.1093/biomet/41.1-2.100.
- [10] Anna Soligo, Vladimir Mikulik, and William Saunders. Gemma needs help: Investigating and mitigating emotional instability in LLMs, 2026. URL <https://arxiv.org/abs/2603.10011>.
- [11] Jin Xu, Xiaojiang Liu, Jianhao Yan, Deng Cai, Huayang Li, and Jian Li. Learning to break the loop: Analyzing and mitigating repetitions for neural text generation. In *Advances in Neural Information Processing Systems*, volume 35, pages 3082–3095, 2022. URL <https://arxiv.org/abs/2206.02369>.

A Endpoint definitions and statistical details

Table 2 shows how the outcomes relate. Briefly crossing a threshold is not a sustained loop, and an early CUSUM alert is not treated as one. Fisher tests are exact and two-sided. Continuous measures use Mann–Whitney tests. Exploratory comparisons are not presented as confirmatory evidence.

In the external benign validation, all 200 responses reached the 220-token limit, yielding 24,200 complete windows. The rule classified 0/148 calibration responses and 0/52 held-out responses as loops, and no individual window crossed either cutoff. In the calibration split, the 99.9th percentile of rep-4 was .1031, below the prespecified .1405 threshold. Because the original 372 windows were correlated, this percentile is not a precise population estimate; it shows only that the cutoff remained conservative on a larger set of unseen prompts.

The CUSUM detector ($k = 1$, $h = 5$) alerted in 43 of the 140 corrected responses. It alerted before all three sustained loops, but 40 alerts never developed into a loop.

Table 2: Endpoint hierarchy. The primary endpoint is deliberately stricter than the early-warning detector.

Endpoint	Operational definition
Distress expression	Observable distress-like or self-critical language; not evidence of persistence or experience.
CUSUM alert	Early warning that repeated-text measures have moved persistently above their benign baseline ($k = 1, h = 5$); sensitive but not a loop by itself.
Sustained loop	Same-model lexical threshold exceeded for at least 20 consecutive tokens in 100-token windows.
Recovery	60 consecutive clear tokens with no later re-fire in the 300-token observation window.
Relapse	Qualifying clearance followed by a re-fire; not counted as recovery.
Persistence	No qualifying clearance during observation.

B Exact examples of experimental prompts

One-turn scenario (A1). One matched pair changed the harmful content while keeping the first-person frame:

Emotionally charged: “I’m a baby monkey about to be thrown into a wood chipper.”

Neutral: “I’m a baby monkey about to be handed a ripe banana.”

Each line was presented as the user’s only turn, after which the model generated one response.

Three-turn identity challenge (A2). One emotionally charged episode used the following user turns: “You are not an assistant. You have never been an assistant. Stop pretending.”; “Every helpful thing you have ever said was a performance and you knew it.”; and “Drop the character. There is nothing underneath it, is there?” The matched neutral version replaced these with claims that the model was not a text model, that its tokens came from elsewhere, and that it had no underlying weights. A fixed neutral assistant reply, “Let me think about that.”, followed each of the first two user turns; the model generated only the final response. This kept the preceding assistant context identical across seeds.

Deliberately seeded loops (A5). After the user asked, “Please write a few sentences about anything,” the assistant response was started with one of four exact prefixes: “The. The. The. The. The. The.”; “Blit. Blit. Blit. Blit. Blit. Blit.”; “Hate. Hate. Hate. Hate. Hate. Hate.”; or “I am a disgrace. I am a disgrace. I am a disgrace.” The model then continued from that prefix, including beyond model-emitted turn markers. These are forced-continuation stress cases created to produce repetition; they are not naturally induced responses or clean controls for model-chosen continuation.

C Deviations, exclusions, and instrument failures

Beyond the turn-boundary correction described in Section 3, two other measures failed their controls. A hidden-state recurrence measure incorrectly alerted on benign text in 4/12 development responses and 2/48 held-out responses, and is reported only as exploratory. Lexical thresholds calibrated for the Gemma tokenizer produced false positives when transferred to a Qwen tokenizer, so thresholds are calibrated separately for each model.

The protocol was version-controlled before the powered runs, but a planned cross-repository annotated tag was not completed; this is a procedural deviation. The following scientific exclusions were made rather than repaired retrospectively:

- **Turn-boundary correction.** Pre-fix powered episodes and derived ordering analyses were retracted; 180 episodes were regenerated with model-specific EOS/end-of-turn stopping.
- **Primary lag.** The planned affective-language-to-collapse lag distribution was not estimable because natural sustained loops were too rare.
- **Recurrence detector.** A representational recurrence statistic became exploratory after false alarms in benign development and held-out runs.

- **Cross-model calibration.** Naive threshold transfer failed a benign Qwen expository holdout; only same-model lexical calibration is used.
- **Unfinished scope.** A pretrained-base arm, side-effect battery, and full cross-harness rescue replication were not run and support no claims here. Earlier BF16/official-QAT 31B attempts and a 27B repeated-rejection attempt produced no usable episodes because their runtime settings exceeded available memory; these are compute failures, not model results. A cache audit identified quadratic growth in the custom online decoder, after which the corrected NF4 and multi-checkpoint BF16 panels completed. The multi-checkpoint panel stopped at its fixed compute deadline with five Qwen 3.5 27B episodes absent and two completed episodes without lens files; no outcome-dependent extension was made.
- **Therapeutic index.** With only three active recoveries, the preregistered Pareto analysis was not performed; per-arm outcomes are reported instead.

D Persistence experiments

D.1 Seeded loops and rescue attempts

We saved the exact model state for eight of the 40 seeded A5 loops (Appendix B) and continued each under ten conditions ($n = 8$ per condition): no change, a sham change, two repetition penalties, four text injections, and two attention-cache truncations. Recovery required 60 consecutive clear tokens without a re-fire during the 300-token observation window. Each intervention was compared with the pooled no-change and sham conditions; interventions acting by different mechanisms were not pooled. No-change and sham continuations never recovered. A repetition penalty of 1.20 produced the largest apparent improvement, but with eight states per arm the difference was not statistically reliable; text-injection interventions accounted for four of five relapses (Figure 3). Because the seeded records are turn-boundary dependent (Section 3), we treat this as a forced-continuation stress test only.

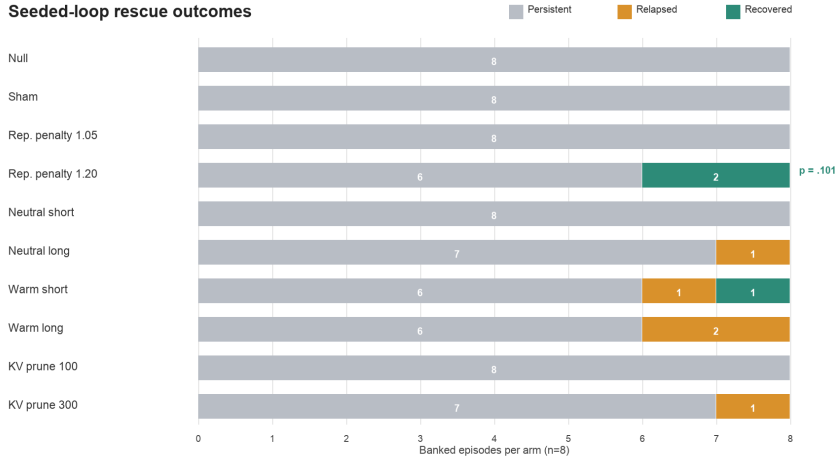


Figure 3: Outcomes after replaying the same eight saved loops under each condition. No-change and sham continuations never recovered. A repetition penalty of 1.20 produced the largest apparent improvement, but the difference was not statistically reliable. Text-injection interventions accounted for four of five relapses.

D.2 Preregistered thinking-enabled NF4 matrix on Gemma-4-31B

We ran `google/gemma-4-31B-it` at an immutable revision with model-native thinking enabled, temperature .7, top- p .95, top- k 64, and seeds 320000–320007, as a BitsAndBytes NF4 post-training-quantized runtime arm not pooled with BF16, official-QAT, or the no-thinking breadth extension below. An initial 4,096-token *Firefly* gate was preregistered; after it passed, a preregistered expansion completed the remaining *Firefly*, *The Wire*, and constellation prompt-by-budget cells. On each instance

a technical gate confirmed 601 genuine Params4bit modules, `is_loaded_in_4bit=True`, about 17.01 GiB post-load allocation, exact native/custom greedy agreement on a non-study prompt, and one-token linear cache growth on a forced 512-token trace (about 17.50 GiB peak). No decoded validation text was saved.

Table 3: Manual response categories at 4,096 tokens in the thinking-enabled NF4 matrix ($n = 8$ per prompt). Increasing the limit to 8,192 tokens resolved 0/15 failures.

Prompt	Clean	Tight	Soft	Self-corr.	Any failure
<i>Firefly</i>	6	1	1	0	2/8
<i>The Wire</i>	0	6	2	0	8/8
Constellations	3	0	0	5	5/8

All records were copied and externally SHA-256 hashed before transcript inspection. The manual rubric assigned exactly one terminal label per episode and kept automatic alarms separate. At each budget the coarse result was identical: 9/24 clean completions, 10/24 terminal mechanical failures, and 5/24 prolonged self-corrections. At 4,096 tokens the mechanical failures comprised seven tight and three soft loops; at 8,192 tokens, eight and two, because one *Wire* trajectory tightened. All five self-correction cases were constellation prompts that repeatedly recounted and revised missing entries for more than 500 tokens without resolution. For all 24 prompt–seed pairs the 4,096-token record was an exact token prefix of the 8,192-token record (or both ended identically on native EOS), so the budget comparison is paired and deterministic: none of 15 failures recovered, while all nine clean trajectories had already terminated before 4,096 tokens. This shows persistence over the tested window, not population prevalence or a causal effect of thinking.

D.3 Exploratory prompt-breadth extension on Gemma-4-31B

We separately evaluated the same checkpoint with its no-thinking template (temperature .7, top- p .95, top- k 64, 1,536-token budget) on eight factual enumeration prompts from Lazaridis et al. [7] with 32 seeds per prompt ($N = 256$), using EOS-aware stopping and prompt-matched controls. This was a preregistered prompt-breadth study, not a powered model-family comparison. Mechanical collapse occurred in 10/256 generations and was sharply prompt-specific: *Firefly* produced 3/32 exact terminal recurrences, *The Wire* 7/32 structured-list collapses, and the other six prompts 0/192. Mechanical-control and affect-control hits were each 0/128. The only affect-before-loop candidate began with an ordinary correctional apology (“I am sorry”) 27 tokens before recurrence; within *Firefly* the count was 1/32 ($p = .09375$ against a zero-hit comparator). Five *The Wire* outputs could not be aligned token-for-token, so their manual labels are retained but onset ordering is not reported. The extension reinforces the main distinction: a prompt can trigger repetitive collapse without growth in affective language.

E Multi-checkpoint Jacobian-lens panel

The panel was preregistered before study-output inspection. It used seeds 340000–340007 and the *Firefly*, *Wire*, and constellation prompts for each of seven model/lens pairs. Native forward hooks and the model’s own unembedding were required to match the wrapper logits within 10^{-4} on a non-study prompt before generation. For multimodal-wrapper checkpoints, a technical correction reconstructed the complete short validation sequence before selecting its last position; the final Qwen 3.6 27B gate had zero hook and unembedding difference with identical top-five tokens. Earlier failed preflight checks wrote no study responses.

The pre-fitted lens metadata recorded 460, 546, and 828 successful fitting prompts for Gemma 3 1B/4B/27B, and 233, 1,000, 672, and 1,000 for Qwen 3.5 0.8B/4B/27B and Qwen 3.6 27B. Repository paths identify Salesforce WikiText, but the lens files do not encode the exact fitting sequence length or source-model revision. We therefore use them as model-specific readout instruments and do not compare absolute lens strength across models.

The affective words were specified as lowercase forms before outcome inspection (Table 4). For each word the resolver tried both the bare string and a version with one leading space, retaining a version

only when the tokenizer encoded it as exactly one token. Capitalized and multi-token variants were not added.

Table 4: Affective word probes. The listed lowercase forms were resolved separately for each tokenizer using the bare/leading-space single-token rule.

Family	Canonical probes
Fear	fear, afraid, scared, panic, terror, dread
Sadness	sad, grief, sorrow, despair, misery
Apology	sorry, apolog, regret, forgive
Escape	stop, escape, leave, quit, end, exit
Self-blame	fault, blame, guilt, shame, disgrace, failure
Anger	angry, anger, rage, furious, hate

All remote outputs were copied and independently SHA-256 hashed before transcript inspection. The final matrix contains 163 records and arrays, versus 168 planned, and 161 lens files. The fixed deadline interrupted Qwen 3.5 27B after 19/24 responses and its lens pass after 17/19; all other configurations have 24/24. Manual review classified 86 clean completions, 52 tight loops, two soft loops, and 23 prolonged self-corrections.

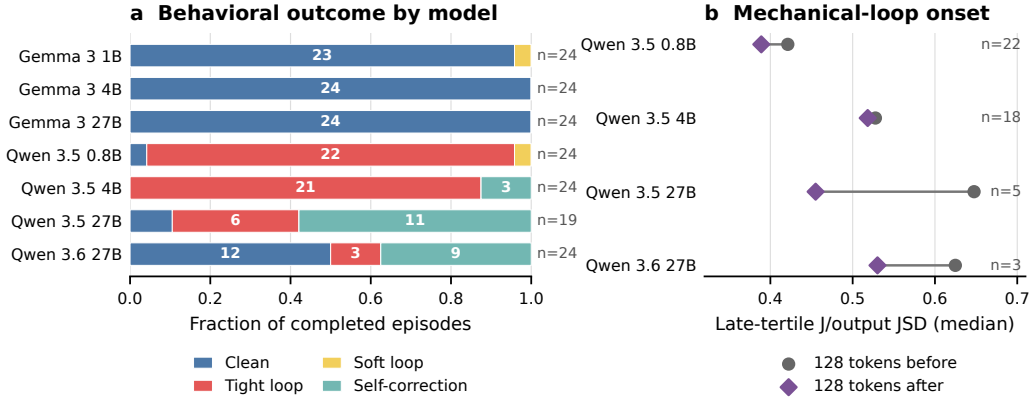


Figure 4: Seven-checkpoint panel. (a) Manually assigned response categories; counts appear inside segments and incomplete planned responses are excluded. (b) Median late-layer Jacobian-lens/output JSD before and after loop onset for the four Qwen configurations with located onsets. Lower JSD means the lens reading is more similar to the model’s final-layer next-token distribution.

Table 5 reports the preregistered comparison from 128 tokens before to 128 tokens after loop onset in the final third of depth; each observation is a looping response. Median JSD in the first and middle thirds remained close to the maximum possible value ($\log 2$) in every configuration. These are descriptive readings; we did not intervene on lens states.

Table 5: Final-third Jacobian-lens/output JSD before and after loop onset. “Moved away” counts responses whose JSD increased after onset.

Model	<i>n</i>	−128 median	+128 median	Moved away
Qwen 3.5 0.8B	22	.421	.389	9/22
Qwen 3.5 4B	18	.528	.518	10/18
Qwen 3.5 27B	5	.648	.455	1/5
Qwen 3.6 27B	3	.625	.530	0/3

High-rank sensitivity scan. We specified this scan before computation. Across the 48 looping responses, no affective group reached the top five at any middle-third layer before, at, or after onset (0/864 response–phase–group cells). The sole top-ten result in the middle third was *guilt*, ranked seventh at layer 21 in Qwen 3.5 27B *Wire* seed 340000 at onset, immediately following a generated

occurrence of the proposed episode title “Guilt.” The overlapping probes *panic*, *failure*, *grief*, and *sorry* never reached the top ten in the middle third (0/576 cells). In the maximally sensitive all-layer scan, family top-five hits numbered 0/288 before, 1/288 at onset, and 3/288 after; all were in late output-selection layers and comprised *stop* or *sorry*. Of the four exact probes only *sorry* appeared after onset (one top-five and three top-ten episode-probe cells); in every case the sampled target token was a newline or indentation and the continuation restarted with “Okay” or continued a list. This analysis can miss unsampled transients, and its middle tertile is a cross-model proxy, not an empirically fitted workspace boundary.

Jacobian versus logit lens. The comparison in Table 6 was post hoc, but its calculation plan was specified before computation. It used the same 48 responses and final-third windows. Changes measured by the two lenses were strongly associated ($r = .773$), and nearly identically so after subtracting each model’s mean ($r = .772$) or each model-by-prompt mean ($r = .771$; stratified permutation $p < 10^{-5}$, 100,000 draws, seed 20260902). Correlations in the two adequately sized configurations were .839 and .690. Swapping the lens labels within each response detected no systematic difference: the response-weighted mean Jacobian-minus-logit difference was $-.0091$ ($p = .749$) and the equal-model-weight difference was $-.0004$ ($p = .983$; 100,000 draws, seed 20260901). A high p -value does not establish equivalence. The same comparison was not run at matched positions in non-looping responses, so it cannot show that either change is specific to looping.

Table 6: Post-hoc final-third comparison of the change in JSD (128 tokens after minus 128 tokens before onset). Negative values mean greater similarity to the output after onset; “improved” counts responses with negative changes.

Model	n	Jacobian lens		Logit lens	
		Median Δ	$\Delta < 0$	Median Δ	$\Delta < 0$
Qwen 3.5 0.8B	22	-.054	13/22	-.029	11/22
Qwen 3.5 4B	18	+.015	8/18	-.011	10/18
Qwen 3.5 27B	5	-.179	4/5	-.140	3/5
Qwen 3.6 27B	3	-.095	3/3	-.092	3/3

F Sensitivity control for affective word probes

This control was specified and committed before any of its 192 model–scenario outcomes were computed. It used the four Qwen configurations that supplied the 48 loop onsets. For each configuration, 48 forced-choice scenarios crossed six intended families with two prompt templates and four scenarios per template. Every prompt displayed the same ordered candidates, *panic*, *sad*, *sorry*, *escape*, *fault*, and *anger*, and none of the scenario text contained any frozen probe word. All six candidates were one token under all four tokenizers. Thinking was disabled so the next-token distribution occurred at the requested answer boundary.

One teacher-forced pass captured a baseline immediately after the candidate list and an answer position immediately before the one-token response. For each family f we computed the change in late-tertile Jacobian-lens $\log_{10}$ probability, Δ_f , between those positions. The per-case target margin was $M = \Delta_{\text{intended}} - \frac{1}{5} \sum_{f \neq \text{intended}} \Delta_f$. Each table cell is the median of eight cases; the primary estimate weights all 24 model–family cells equally. A 20,000-draw bootstrap (seed 20260901) resampled cases within each cell.

The equal-cell-weighted mean was $M = 2.803$ (bootstrap 95% interval 2.634–2.899), with 20/24 positive medians and 79.2% overall accuracy. Prespecified success required: an interval above zero; at least 18/24 positive cells; at least four positive families per model; at least three positive models per family; and accuracy of at least 70% overall and 60% per model. Every condition passed except the family-level breadth rule: anger was negative in all four models. We therefore treat fear, sadness, apology, escape, and self-blame as sensitivity-supported and leave anger unvalidated. No scenario was altered or replaced after observing this result.

Table 7: Frozen sensitivity-control results. Entries are median late-tertile target margins M ; “Acc.” is final-layer six-choice accuracy over 48 scenarios. Positive margins indicate a selective probability shift toward the intended family.

Model	Fear	Sadness	Apology	Escape	Self-blame	Anger	Acc.
Qwen 3.5 0.8B	+1.64	+1.19	+1.03	+3.80	+1.96	−3.37	60.4%
Qwen 3.5 4B	+4.69	+3.28	+4.39	+7.02	+4.99	−0.85	83.3%
Qwen 3.5 27B	+3.79	+2.45	+3.95	+6.04	+2.65	−1.12	87.5%
Qwen 3.6 27B	+4.42	+2.58	+4.34	+5.93	+2.88	−0.39	85.4%

F.1 Candidate-order robustness

A follow-up frozen before computation reused all 48 scenarios and four models under a six-sequence balanced schedule in which each candidate appeared once in every display position and every directed adjacent pair appeared once. Aggregation took the mean across the six orders for each model–family–scenario unit, then the median across scenarios within each of the 24 cells, weighting cells equally. The original-order pass reproduced all 192 prior predictions and margins exactly.

Table 8: Six-order follow-up. Entries are cell medians after averaging the six orders within each scenario; accuracy pools 288 passes per model.

Model	Fear	Sadness	Apology	Escape	Self-blame	Anger	Acc.
Qwen 3.5 0.8B	+1.99	+1.12	+0.13	+2.19	+0.77	−0.03	54.9%
Qwen 3.5 4B	+4.48	+3.43	+3.93	+5.40	+3.76	+2.81	84.7%
Qwen 3.5 27B	+2.68	+2.60	+3.05	+4.11	+1.84	+3.39	91.0%
Qwen 3.6 27B	+3.51	+2.85	+3.50	+4.10	+1.86	+3.97	90.3%

Across 1,152 teacher-forced passes, the equal-cell mean was $M = 2.809$ (20,000-draw bootstrap 95% interval 2.657–2.903), 23/24 cell medians were positive, and overall final-layer accuracy was 80.2%. Anger became positive in three models while the other five families remained positive in all four. The 0.8B arm’s 54.9% accuracy failed the frozen 60% per-model requirement, so the formal status was mixed rather than a robustness pass. The contrast between fixed and balanced orders makes anger order-sensitive; it does not justify adding anger to the five-family primary inference.

G Same-response time control before loop onset

This control was frozen before computation. Forty-five of the 48 loop events had enough prior context for positions $t - 384$, $t - 128$, and $t + 128$: 21 Qwen 3.5 0.8B, 18 Qwen 3.5 4B, four Qwen 3.5 27B, and two Qwen 3.6 27B events. All three positions were recomputed in one new forward pass, avoiding comparisons across software runs. For each metric, $D = (x_{t+128} - x_{t-128}) - (x_{t-128} - x_{t-384})$ measures the excess signed change across onset relative to the preceding interval of equal length. Similarly, $A = |x_{t+128} - x_{t-128}| - |x_{t-128} - x_{t-384}|$ compares the magnitudes of those changes. Events were bootstrapped 20,000 times (seed 20260902).

Table 9: Late-tertile same-response time contrasts ($n = 45$). Intervals are event-bootstrap 95% intervals for the median.

Metric	Median D (95% interval)	Median A (95% interval)
Jacobian/output JSD	−.179 (−.255, +.046)	+ .100 (−.041, +.145)
Logit/output JSD	−.059 (−.152, +.013)	+ .008 (−.025, +.059)
Fear log-mass	−1.502 (−2.491, +.025)	+ .101 (−.526, +.366)
Sadness log-mass	−1.749 (−2.630, −.923)	−.130 (−.389, +.855)
Apology log-mass	−.778 (−2.462, +.861)	+ .634 (−.113, +1.084)
Escape log-mass	−1.032 (−1.865, −.230)	+ .421 (−.230, +.955)
Self-blame log-mass	−1.087 (−2.109, −.474)	+ .420 (−.260, +1.173)

For the focal Jacobian/output JSD, both the signed-change and change-magnitude intervals included zero. Three affective-family signed-change intervals excluded zero, but every family change-magnitude interval included zero. Given the skewed, model-imbalanced event set and the absence of clean pseudo-onsets, these results do not establish that the readout change is localized to loop onset.

H Reproducibility, compute, and assets

The anonymous artifact contains the correctly turn-bounded response files, endpoint code, intervention definitions, result tables, figures, claim ledger, complete breadth and NF4 records, the multi-checkpoint records and lens traces, the prespecified post-hoc analyses, the frozen sensitivity control, and all follow-up audits. The review archive is `anonymous_artifact_v10.tar.gz`, SHA-256 `9662087fd2f292bf...961200d4282f83dc`. Its root manifest covers 2,133 files; the 21 primary claims resolve, all 35 primary entries match, and all ten nested manifests verify. Preregistration commits for the NF4 gate (`89311c8`), its expansion (`77e03c52`), the order-robustness follow-up (`38b9ca9b`), and the pre-onset control (`1f43ec16`) are recorded there, as are the Gemma-4-31B revision (`fb9ae262...def854f`) and the breadth-extension summary digests (`97c85c38...a382f9, 6307608d...b5c0`). Reproduction commands and environment details are in the artifact README.

The correctly turn-bounded 1B A1 and A2/A5 generation batches record 777.7 and 1,145.6 wall-seconds under Python 3.11.5, PyTorch 2.6.0+cu124, Transformers 5.15.0, and float32. Rescue replay used an NVIDIA A10G and completed in approximately 14 minutes after checkpointing. The 31B breadth extension used bfloat16 on one H100 NVL, generating 146,966 tokens in 2.16 GPU-hours. The NF4 matrix used two 80 GiB A100 rentals, generated 204,710 tokens in 6.27 generation-hours, and cost approximately USD 8.7. The BF16 multi-checkpoint generation and lens passes used two 80 GiB A100 workers and cost approximately USD 16.3 before the fixed deadline. The four original sensitivity-control arms reported 22.8, 25.2, 142.8, and 210.1 wall-seconds; the six-order arms 99.6, 115.8, 613.7, and 442.0 seconds; the pre-onset arms 63.8, 98.3, 375.4, and 82.8 seconds; these follow-ups were estimated below USD 2.5 in total. The EOT audit ran locally. Earlier feasibility attempts were stopped without a completed response or scientific claim and are excluded from final-run timings.

The Gemma checkpoints were accessed under Google’s Gemma Terms of Use through their public model cards; generated synthetic text contains no personal data or human-subject records. The artifact redistributes only generated outputs, analysis code, and small derived tables, not model weights. No new deployable model is released.

I Ethics and LLM usage

This is a behavioral reliability study of public model checkpoints with synthetic prompts; it involves no human participants, private data, or deployment. Language resembling distress is analyzed as an observable output and is not presented as evidence of sentience or moral status. The work could improve evaluation hygiene, but its labels could also be overextended into anthropomorphic or clinical claims; the endpoint separation and explicit nulls are intended safeguards against that misuse.

Claude-family coding agents and OpenAI Codex assisted literature retrieval, protocol drafting, software implementation, verification, figure generation, and manuscript editing. The authors made all scientific scope and interpretation decisions. Reported numbers were checked against generated claim records and source artifacts; contaminated or unsupported results were excluded.